

Fiat500_VehicleSelection

June 12, 2023

DATE:-6-6-2023 ____FIAT500_VEHICLE SELECTION USING LINEAR , LASSO ,
RIDGE

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
[ ]: df=pd.read_csv(r"/content/fiat500_VehicleSelection_Dataset (1).csv")
df
```

```
[ ]:
      ID  model  engine_power  age_in_days      km  previous_owners  \
0      1  lounge           51          882   25000                1
1      2    pop           51         1186   32500                1
2      3   sport           74         4658  142228                1
3      4  lounge           51         2739  160000                1
4      5    pop           73         3074  106880                1
...    ...    ...    ...    ...    ...    ...
1533  1534   sport           51         3712  115280                1
1534  1535  lounge           74         3835  112000                1
1535  1536    pop           51         2223   60457                1
1536  1537  lounge           51         2557   80750                1
1537  1538    pop           51         1766   54276                1
```

```
      lat      lon  price
0  44.907242  8.611560   8900
1  45.666359  12.241890   8800
2  45.503300  11.417840   4200
3  40.633171  17.634609   6000
4  41.903221  12.495650   5700
...    ...    ...    ...
1533  45.069679  7.704920   5200
1534  45.845692  8.666870   4600
1535  45.481541  9.413480   7500
1536  45.000702  7.682270   5990
```

```
1537  40.323410  17.568270  7900
```

```
[1538 rows x 9 columns]
```

```
[ ]: df.head()
```

```
[ ]:      ID  model  engine_power  age_in_days      km  previous_owners      lat  \
0     1  lounge           51          882   25000              1  44.907242
1     2    pop           51         1186   32500              1  45.666359
2     3  sport           74         4658  142228              1  45.503300
3     4  lounge           51         2739  160000              1  40.633171
4     5    pop           73         3074  106880              1  41.903221

      lon  price
0   8.611560   8900
1  12.241890   8800
2  11.417840   4200
3  17.634609   6000
4  12.495650   5700
```

```
[ ]: df.tail()
```

```
[ ]:      ID  model  engine_power  age_in_days      km  previous_owners  \
1533  1534  sport           51         3712  115280              1
1534  1535  lounge           74         3835  112000              1
1535  1536    pop           51         2223   60457              1
1536  1537  lounge           51         2557   80750              1
1537  1538    pop           51         1766   54276              1

      lat      lon  price
1533  45.069679  7.70492   5200
1534  45.845692  8.66687   4600
1535  45.481541  9.41348   7500
1536  45.000702  7.68227   5990
1537  40.323410 17.56827   7900
```

```
[ ]: df.shape
```

```
[ ]: (1538, 9)
```

```
[ ]: df.describe
```

```
[ ]: <bound method NDFrame.describe of      ID  model  engine_power  age_in_days
km  previous_owners  \
0      1  lounge           51          882   25000              1
1      2    pop           51         1186   32500              1
2      3  sport           74         4658  142228              1
```

3	4	lounge	51	2739	160000	1
4	5	pop	73	3074	106880	1
...
1533	1534	sport	51	3712	115280	1
1534	1535	lounge	74	3835	112000	1
1535	1536	pop	51	2223	60457	1
1536	1537	lounge	51	2557	80750	1
1537	1538	pop	51	1766	54276	1

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]>

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model           1538 non-null   object
2   engine_power    1538 non-null   int64
3   age_in_days     1538 non-null   int64
4   km              1538 non-null   int64
5   previous_owners 1538 non-null   int64
6   lat             1538 non-null   float64
7   lon             1538 non-null   float64
8   price           1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

```
[ ]: df.isna().any()
```


```
[ ]: ID           False
      model        False
```

```
engine_power      False
age_in_days       False
km                False
previous_owners   False
lat               False
lon               False
price             False
dtype: bool
```

```
[ ]: df.isna().any()
```

```
[ ]: ID          False
      model       False
      engine_power False
      age_in_days False
      km          False
      previous_owners False
      lat         False
      lon         False
      price       False
      dtype: bool
```

```
[ ]: sns.lmplot(x='km',y='price',data=df,order=2,ci=None)
      plt.show()
```




```
[ ]: x=np.array(df['km']).reshape(-1,1)
     y=np.array(df['price']).reshape(-1,1)
```

```
[ ]: df.dropna(inplace=True)
```


```
[ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
     #splitting data into train and test
     regr=LinearRegression()
     regr.fit(x_train,y_train)
     print(regr.score(x_test,y_test))
```

0.7415041964836075

```
[ ]: y_pred=regr.predict(x_test)
     plt.scatter(x_test,y_test,color='b')
     plt.plot(x_test,y_pred,color='k')
     plt.show()
```



```
[ ]: df500=df[:][:500]
sns.lmplot(x="km",y="price",data=df500,order=1,ci=None)
plt.show()
```



```
[ ]: from sklearn.linear_model import LinearRegression
      from sklearn.metrics import r2_score
```

```
[ ]: #train model
      model=LinearRegression()
      model.fit(x_train,y_train)
      #Evaluation the model on the test set
      y_pred=model.predict(x_test)
      r2=r2_score(y_test,y_pred)
      print("R2 score:",r2)
```

R2 score: 0.7415041964836075


```
[ ]: from sklearn.linear_model import Lasso,Ridge
      from sklearn.preprocessing import StandardScaler
```

```
[ ]: plt.figure(figsize = (10, 10))
      sns.heatmap(df500.corr(), annot = True)
```

```
plt.show()
```

<ipython-input-19-a0241e6b6c90>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df500.corr(), annot = True)
```




```
[ ]: features=df.columns[0:1]
      target=df.columns[-1]
```

```
[ ]: X = df[features].values
      y = df[target].values
      #split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
      ↪random_state=17)
      print("The dimension of X_train is {}".format(X_train.shape))
      print("The dimension of X_test is {}".format(X_test.shape))
      #Scale features
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

The dimension of X_train is (1153, 1)

The dimension of X_test is (385, 1)

```
[ ]: lr = LinearRegression()
      #Fit model
      lr.fit(X_train, y_train)
      #predict
      #prediction = lr.predict(X_test)
      #actual
      actual = y_test
      train_score_lr = lr.score(X_train, y_train)
      test_score_lr = lr.score(X_test, y_test)
      print("\nLinear Regression Model:\n")
      print("The train score for lr model is {}".format(train_score_lr))
      print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.00310286926477088

The test score for lr model is -0.008405634316406507

```
[ ]: #Ridge Regression Model
      ridgeReg = Ridge(alpha=10)
      ridgeReg.fit(X_train,y_train)
      #train and test scorefor ridge regression
      train_score_ridge = ridgeReg.score(X_train, y_train)
      test_score_ridge = ridgeReg.score(X_test, y_test)
      print("\nRidge Model:\n")
      print("The train score for ridge model is {}".format(train_score_ridge))
      print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.0031026398591535997


The test score for ridge model is -0.008307809466001403

```
[ ]: plt.figure(figsize=(10,10))
```

```
[ ]: <Figure size 1000x1000 with 0 Axes>
```

```
<Figure size 1000x1000 with 0 Axes>
```

```
[27]: plt.plot(features,ridgeReg.coef_,alpha=0.  
↳7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge;  
↳$\alpha=10$',zorder=7)  
plt.plot(features,lr.coef_,alpha=0.  
↳4,linestyle='none',marker="o",markersize=7,color='green',label='LinearRegression')  
plt.xticks(rotation=90)  
plt.legend()  
plt.show()
```



```
[31]: #Lasso regression model  
print("\nLasso Model: \n")
```

```

lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))

```

Lasso Model:

The train score for ls model is 0.003075838461310987


The test score for ls model is -0.007367578602064606

```

[32]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "
      ↪"bar")

```

[32]: <Axes: >



```

[33]: #Using the linear CV model
      from sklearn.linear_model import LassoCV
      #Lasso Cross validation

```

```

lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).
    ↪fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))


```

0.0031025989567363688
-0.008299466692577973

```

[34]: #plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.
    ↪7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;␣
    ↪$\alpha=10$',zorder=7)
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.
    ↪5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;␣
    ↪$\alpha = \text{grid}$')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.
    ↪4,linestyle='none',marker='o',markersize=7,color='green',label='Linear␣
    ↪Regression')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()

```



```
[35]: #Using the linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train,
    ↪ y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train,
    ↪ y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test,
    ↪ y_test)))
```

The train score for ridge model is 0.0031026398591535997
 The train score for ridge model is -0.008307809466002958

```
[36]: from sklearn.linear_model import ElasticNet
      regr=ElasticNet()
      regr.fit(X,y)
      print(regr.coef_)
      print(regr.intercept_)
```

```
[0.12455754]
8480.156871173602
```

```
[39]: y_pred_elastic=regr.predict(X_train)
```

```
[40]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
      print(mean_squared_error)
```

```
3708273.194830543
```